## Peer Response

by Jaafar El Komati - Wednesday, 1 October 2025, 9:08 AM
Abdulrahman, I really liked how you clearly carried forward risks and practical methods. Talking about the costs hidden in data labeling, I got the most gone. Only the "large" moral debate is easy to focus on, such as prejudice or misinformation, but the interaction is made incomplete to see the labor behind the AI. Recent reports have highlighted how the anotation work is often outsourced in Global South, which raises questions of digital exploitation (Gray & Suri, 2019).

Your emphasis on data consent and traceability also takes place on time. This year is still coming out this year with cases against AI companies, it is clear that " managing data well " is not just a technical reform - it has become a legal and cultural requirement (Vincent, 2023). I think your suggestion to embed watermark for Provence is particularly interesting. It can help create accountability, although it also raises questions about enforcement in the open-source ecosystem.

I strongly agree with your call to keep humans in the loop. Transparency and human inspection feel like a bridge between innovation and trust. Without it, we risk the systems that look efficient but take silent loss.

### References


- Gray, M.L. & Suri, S. (2019). *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass*. Houghton Mifflin Harcourt.
- Vincent, J. (2023). Artists file copyright lawsuits against AI image generators. *The Verge*. Available at: https://www.theverge.com

## Peer Response

by Shaikah Salim Mohammed Alkhaayyal Alharthi - Sunday, 5 October 2025, 6:59 PM

Your assessment addresses the many facets of the ethical dimension of Deep Learning in a coherent, pragmatic, and structured way. I especially applaud the way you have connected the misuse of data to legal and moral responsibility. As pointed out by Birhane and Prabhu (2021), the use of datasets without permission constitutes a violation of rights and a declining trust in AI. This declining trust, in turn, weakens the public's belief in the innovation of technology and its long-term trustworthiness.

Bias, misinformation, and the large-scale amplifying of harmful stereotypes as narrated by Bender et al. (2021) and Bommasani et al. (2022) contain fictitious narratives poses a significant threat. This is why there is an obvious need to conduct strict and thorough testing of systems before and after deployment.

I appreciated the emphasis you place on human oversight. As you state, having humans in the loop provides accountability and ensures automated systems don't make ethical decisions autonomously. Constructing mechanisms to embed watermarks and source tracking are excellent ways to promote transparency.

Including safeguards, such as governance of datasets and the auditing of bias as you have suggested, provides a dignified and trustworthy basis upon which responsible AI may be built.

## Peer response

by Ahmed Ayman Abdelrahman - Monday, 6 October 2025, 4:09 PM

Thank you, Abdelrahman, for an insightful and well-structured post. I particularly liked how you organised your ideas into “Key Dangers” and “Ways Forward” — it provides a clear framework for evaluating the ethical landscape of deep learning. Your discussion on data consent and bias (Birhane & Prabhu, 2021; Bender et al., 2021) highlights two of the most pressing challenges in today’s generative-AI systems. These issues illustrate that technological progress often outpaces our social and legal capacity to govern it.

I also found your point about the “hidden costs” of data labelling especially important. As Gray and Suri (2019) note, the invisible workforce behind AI systems often faces poor working conditions and minimal recognition. Addressing AI ethics must therefore extend beyond algorithmic fairness to include labour ethics and supply-chain transparency.

Your proposed solutions — such as dataset tracking, watermarking, and human oversight — offer a solid foundation for responsible AI development. I would add that algorithmic impact assessments (Reisman et al., 2018) could complement these strategies by systematically evaluating potential harms before deployment.

Overall, your post captures the balance between innovation and accountability effectively. Do you believe that implementing such safeguards should be mandated through regulation, or left to voluntary corporate governance frameworks?

## Peer Response

by Saleh Almarzooqi - Thursday, 9 October 2025, 1:42 AM

I would like to thank you for presenting such a succinct yet detailed list of risks and potential security measures against deep learning systems. I found it very helpful that your post systematically linked the ethical issues to practical action, instead of discussing them abstractly.

Your focus on data consent and provenance, in particular, I find especially appropriate. There are no clear records of the origin of training data, so that bias, intellectual property, or privacy risks cannot be evaluated. To build on your argument, according to Brajovic, Danilo, et al. (2023), there may be the introduction of so-called model cards and data cards as standard documentation to make end-users and regulators understand how a model was constructed and tested.

Embedding watermarks and tags is also an important solution to address the misinformation that you suggested. A recent study by Cao (2025) indicates that strong watermarking, along with open provenance specifications, can assist websites and news reporters in locating AI-generated material before it goes viral.

And lastly, I like your observation of the keeping of human beings in a circle. Besides review processes, organisations might introduce obvious red lines in areas where AI cannot be employed without direct supervision, such as in medical or legal decision-making (Jago, Robert, et al. 2021). This would ease the accountability and trust, and still enjoy the benefits of automation.

Generally speaking, your post emphasises that the development of AI responsibly requires a combination of technical, legal, and organisational efforts instead of a single solution.

References:

Brajovic, D., Renner, N., Goebels, V.P., Wagner, P., Fresz, B., Biller, M., Klaeb, M., Kutz, J., Neuhuettler, J. and Huber, M.F., 2023. Model reporting for certifiable AI: A proposal for merging EU regulation into AI development. arXiv preprint arXiv:2307.11525.

Cao, L., 2025. A Practical Synthesis of Detecting AI-Generated Textual, Visual, and Audio Content. arXiv preprint arXiv:2504.02898.

Jago, R., van der Gaag, A., Stathis, K., Petej, I., Lertvittayakumjorn, P., Krishnamurthy, Y., Gao, Y., Silva, J.C., Webster, M., Gallagher, A., and Austin, Z., 2021. Use of artificial intelligence in regulatory decision-making. Journal of Nursing Regulation, 12(3), pp.11-19.