

Peer Response

by [Abdulrahman Alhashmi](#) - Thursday, 2 October 2025, 9:24 PM

Ahmed, you describe the key risks clearly. I agree with your points on false information, intellectual property, bias, and accountability. To move from identifying issues to practical steps, I suggest a straightforward framework that teams can apply at once.

- Track data and permissions. Follow “Datasheets for Datasets” to note where data comes from, its licences, who it covers, and any risks. Drop data with fuzzy rights and cut storage time (Gebru et al., 2021).
- Share model details and flaws. Release Model Cards that cover planned uses, risks from misuse, test results by group, and steps to handle failures (Mitchell et al., 2019).
- Add origin tags. Link Content Credentials or similar markers to content. Combine them with stats-based watermarks for text and pictures. This builds trust and aids checks (Kietzmann et al., 2020).
- Check often. Require outside audits before launch, regular follow-ups, and reviews after issues. Link them to firm risk levels. See audit results as must-fix blocks, not tips (Raji et al., 2020).
- Involve people in key decisions. Demand two-person checks for jobs, loans, health care, and law tasks. Log all actions for clear tracking (Bender et al., 2021).

These actions won't harm new ideas; they guide them with care. Growth alone won't solve deep problems. Solid records, open reports, and real checks turn your governance plan into something teams can test and use.

References

- Bender, E.M., Gebru, T., McMillan-Major, A. and Shmitchell, S. (2021) On the dangers of stochastic parrots: Can language models be too big? Proceedings of FAccT '21, pp. 610–623.
- Gebru, T. et al. (2021) Datasheets for datasets. Communications of the ACM, 64(12), pp. 86–92.
- Kietzmann, J., Lee, L., McCarthy, I.P. and Kietzmann, T. (2020) Deepfakes: Trick or treat? Business Horizons, 63(2), pp. 135–146.
- Raji, I.D. et al. (2020) Closing the AI accountability gap: Defining auditing for AI. Proceedings of FAccT '20, pp. 33–44.

Peer Response

by [Abdulrahman Alhashmi](#) - Saturday, 4 October 2025, 3:01 AM

Your post highlights real risks in authorship, bias, and misuse. I agree with the core of your arguments. I add two ideas: separate model skills from rule-setting power. Assess uses by evidence of reduced harm, not just intentions.

For copyright and ownership, data sources form the main issue. Lacking solid records of data origins and consent methods, claims to original work remain weak. Dataset inventories and source-tracking standards allow large-scale checks of origins and rights (Gebru et al., 2018; C2PA, 2023). Such transparency supports fair agreements and targeted deletions.

The bias concerns you note are valid. Bias stems from training data and tool applications. Connect risks to the specific job. An art tool encounters different threats than a hiring tool. A straightforward approach—identify risks, test them, refine—grounds discussions in actual outcomes (Bender et al., 2021; NIST, 2023).

On misuse, open-ended outputs create spam, deepfakes, and targeted falsehoods. Layered protections work well: traceable content, bans on harmful actions, and post-use reviews (Weidinger et al., 2022; Kirchenbauer et al., 2023).

Practical actions: • Insist on data logs, approved sources, and deletion options before training or updates. • Mandate pre-release audits, with tests for bias, safety, and misuse linked to job-specific threats. • Embed safeguards early: usage caps, blocks on dangerous queries, and activity logs. • Apply source labels (C2PA) and technical watermarks to support attribution and compliance. • Commit to external audits and share repeatable test summaries.

These steps do not erase all dangers. They clarify accountability—and enable responses.

References

Bender, E.M., Gebru, T., McMillan-Major, A. and Shmitchell, S. (2021) 'On the Dangers of Stochastic Parrots', Proceedings of FAccT '21. Available at: <https://dl.acm.org/doi/10.1145/3442188.3445922> (Accessed: 4 October 2025).

Coalition for Content Provenance and Authenticity (C2PA) (2023) Technical Specification. Available at: <https://c2pa.org> (Accessed: 4 October 2025).

Gebru, T. et al. (2018) 'Datasheets for Datasets', arXiv:1803.09010. Available at: <https://arxiv.org/abs/1803.09010> (Accessed: 4 October 2025).

Kirchenbauer, J. et al. (2023) 'A Watermark for Large Language Models', arXiv:2301.10226. Available at: <https://arxiv.org/abs/2301.10226> (Accessed: 4 October 2025).

NIST (2023) AI Risk Management Framework 1.0. Available at: <https://www.nist.gov/itl/ai-risk-management-framework> (Accessed: 4 October 2025).

Weidinger, L. et al. (2022) 'Taxonomy of Risks Posed by Language Models', arXiv:2112.04359. Available at: <https://arxiv.org/abs/2112.04359> (Accessed: 4 October 2025).

Peer Response

by [Abdulrahman Alhashmi](#) - Sunday, 5 October 2025, 1:33 PM

Saleh, your worries point to main risks. Moral issues stand out as pressing and plain. I offer three solid protections.

Start by seeing AI generators as tools with two sides. Look at possible damage, not only errors. Apply risk scans, slow launches, and misuse strategies. They reduce chances of broad lies from false content (Brundage et al., 2018; Chesney and Citron, 2019).

Next, tighten data guidelines. Rely on checked, shared sources. Remove private facts and scan for data leaks ahead. Models tend to repeat learned parts, so run privacy audits and group trials to halt release. Record problems and solve them before going live (Carlini et al., 2021; Weidinger et al., 2021).

Next, add duty to daily jobs. Outside checks, data records, and error follows build real trust. Mix source tags (useful yet narrow) with online guidelines: warnings for bold claims, proof connections, and limits on high-risk parts. Link product drops to expert checks. Track live use with plain halt steps (Raji et al., 2020).

Such steps steer growth, without halting it. Your points on bias, rights, and privacy form a complete defence. Safety starts in planning, flows to trials, and holds in oversight. This loop plan, test, assess, track guards people, holds trust, and yields true benefits.

References

Brundage, M. et al. (2018) 'The Malicious Use of Artificial Intelligence: Forecasting, Prevention, and Mitigation', arXiv.

Carlini, N. et al. (2021) 'Extracting Training Data from Large Language Models', Proceedings of USENIX Security Symposium.

Chesney, R. and Citron, D.K. (2019) 'Deep Fakes: A Looming Challenge for Privacy, Democracy, and National Security', California Law Review, 107(6), pp. 1753–1819.

Raji, I.D. et al. (2020) 'Closing the AI Accountability Gap: Defining Auditing and Oversight for Algorithmic Systems', Proceedings of the ACM FAccT Conference, pp. 33–44.

Weidinger, L. et al. (2021) 'Ethical and Social Risks of Harm from Language Models', arXiv.