

Summary Post

by [Abdulrahman Alhashmi](#) - Tuesday, 7 October 2025, 7:09 AM

We reached a balanced view in our talk. Deep learning brings real gains. Yet these depend on consent, fairness, and clear duties. Our blog posts named five key risks. They also list steps to fight them.

Key risks. Data taken without okay hurts artists and breaks trust from the public (Birhane and Prabhu, 2021; Vincent, 2023). Bias baked in leads to stereotypes and errors that hit groups in unfair ways (Bender et al., 2021). Fake results push bad info far and wide (Bommasani et al., 2022). Leaks from data spill private facts (Carlini et al., 2021). Hidden labour links AI rise to jobs that can vanish fast (Gray and Suri, 2019).

Steps to take • Data rules: track sources and get consent first with "datasheets" (Gebru et al., 2018). • Tough checks: test for bias, measure accuracy, and run safety drills before and after launch (Bommasani et al., 2022). • Openness: add watermarks and source tags for easy tracking. • Human checks: assign people to key choices with clear ways to raise issues. • Reviews of effects for systems that touch the public (Reisman et al., 2018).

Others built on the ideas. Jaafar focused on fair work rules and laws that push for consent. Shaikah stressed trust via tests and human checks. Ahmed pushed for reviews of algorithm effects to boost planning ahead.

My view: create basic laws on consent, safety, and openness; link them to solid engineering like small tests, public records, and ongoing watches. Make systems that help. Cut down harm. Log choices and results.

Short term steps: gather legal data sets; set standard data notes; support labelers' well being; share model and system reports; use event logs; prepare for fixes and resets. Open-source groups need custom plans so tracking and watermarks aid research without stopping it, yet still hold people accountable.

These steps match the goals: build with clear risk guesses, launch with safeguards, and check with proof. Most of all, weave ethics into daily tasks, not just as an add-on or rule following, ever.

References

Bender, E.M. et al. (2021) 'On the Dangers of Stochastic Parrots'. In: FAccT '21. New York: ACM.

Birhane, A. and Prabhu, V.U. (2021) 'Large image datasets: A pyrrhic victory?' WACV Workshops. IEEE.

Bommasani, R. et al. (2022) 'On the Opportunities and Risks of Foundation Models'. arXiv:2108.07258.

Carlini, N. et al. (2021) 'Extracting Training Data from Large Language Models'. USENIX Security.

Gebru, T. et al. (2018) 'Datasheets for Datasets'. arXiv:1803.09010.

Gray, M.L. and Suri, S. (2019) Ghost Work. Boston: HMH.

Reisman, D. et al. (2018) Algorithmic Impact Assessments. AI Now Institute.

Vincent, J. (2023) 'Artists file copyright lawsuits against AI image generators'. The Verge.